

Greedy and Local Ratio Algorithms in the MapReduce Model

Nick Harvey¹ Chris Liaw¹ **Paul Liu**²



¹University of British Columbia

²Stanford University


MapReduce at a glance [Vassilvitskii, 2012]



Unordered Data




MapReduce at a glance [Vassilvitskii, 2012]



MapReduce at a glance [Vassilvitskii, 2012]



Unordered Data



Map


Shuffle



MapReduce at a glance [Vassilvitskii, 2012]




Unordered Data




Map

Shuffle



Reduce

Unordered Data




Data:

- ▶ Represented as $\langle \text{Key}, \text{Value} \rangle$ pairs


Operations:

- ▶ Map: $\langle \text{Key}, \text{Value} \rangle \rightarrow \text{List}(\langle \text{Key}, \text{Value} \rangle)$
- ▶ Shuffle: Aggregate all pairs with the same key
- ▶ Reduce: $\langle \text{Key}, \text{List}(\text{Value}) \rangle \rightarrow \langle \text{Key}, \text{List}(\text{Value}) \rangle$




- ▶ n number of vertices
- ▶ $m = n^{1+c}$ input size (number of edges)
- ▶ $n^{1+\mu}$ memory on each machine, $0 < \mu \leq c$
- ▶ $n^{c-\mu}$ machines

Efficiency: algorithms that run in $O(1)$ rounds (often $O(c/\mu)$).



(a) arXiv



(b) Patents

Figure: $|E| \approx n^{1.69}$ (left) and $n^{1.66}$ (right)

- ▶ Filtering: unweighted vertex cover, matching, edge cover [Lattanzi et al., 2011]
- ▶ Rounding and linear programming for weighted matching [Lattanzi et al., 2011, Crouch and Stubbs, 2014, Grigorescu et al., 2016, Ahn and Guha, 2015]

- ▶ Filtering: unweighted vertex cover, matching, edge cover [Lattanzi et al., 2011]
- ▶ Rounding and linear programming for weighted matching [Lattanzi et al., 2011, Crouch and Stubbs, 2014, Grigorescu et al., 2016, Ahn and Guha, 2015]


Our paper:

- ▶ Randomized local ratio (weighted matching, vertex cover, etc.)
- ▶ Hungry-greedy (maximal indep. set, maximal clique, etc.)
- ▶ Additional results on colouring problems


Randomized Local Ratio

- ▶ Approximation technique for NP-hard optimization problems
- ▶ Initially used in the 80s for finding low cost vertex covers [Bar-Yehuda and Even, 1985]
- ▶ Recently used by [Paz and Schwartzman, 2017] to find weighted matchings in the streaming model


A maximum weight matching is a set $M \subseteq E$ s.t. no two $e \in M$ shares an endpoint and the sum of weights in M is maximized.




A maximum weight matching is a set $M \subseteq E$ s.t. no two $e \in M$ shares an endpoint and the sum of weights in M is maximized.



Select **any** edge e with positive weight. Reduce its weight from itself and its neighboring edges. Push e onto a stack. Repeat this procedure until no positive weight edges remain. Pop edges off the stack, adding greedily to the matching.




Select **any** edge e with positive weight. Reduce its weight from itself and its neighboring edges. Push e onto a stack. Repeat this procedure until no positive weight edges remain. Pop edges off the stack, adding greedily to the matching.




S:


Select **any** edge e with positive weight. Reduce its weight from itself and its neighboring edges. Push e onto a stack. Repeat this procedure until no positive weight edges remain. Pop edges off the stack, adding greedily to the matching.




S: 

Select **any** edge e with positive weight. Reduce its weight from itself and its neighboring edges. Push e onto a stack. Repeat this procedure until no positive weight edges remain. Pop edges off the stack, adding greedily to the matching.




S: 


Select **any** edge e with positive weight. Reduce its weight from itself and its neighboring edges. Push e onto a stack. Repeat this procedure until no positive weight edges remain. Pop edges off the stack, adding greedily to the matching.



Select **any** edge e with positive weight. Reduce its weight from itself and its neighboring edges. Push e onto a stack. Repeat this procedure until no positive weight edges remain. Pop edges off the stack, adding greedily to the matching.




Select **any** edge e with positive weight. Reduce its weight from itself and its neighboring edges. Push e onto a stack. Repeat this procedure until no positive weight edges remain. Pop edges off the stack, adding greedily to the matching.



S:




Select **any** edge e with positive weight. Reduce its weight from itself and its neighboring edges. Push e onto a stack. Repeat this procedure until no positive weight edges remain. Pop edges off the stack, adding greedily to the matching.



S:




Select **any** edge e with positive weight. Reduce its weight from itself and its neighboring edges. Push e onto a stack. Repeat this procedure until no positive weight edges remain. Pop edges off the stack, adding greedily to the matching.



S:




Select **any** edge e with positive weight. Reduce its weight from itself and its neighboring edges. Push e onto a stack. Repeat this procedure until no positive weight edges remain. Pop edges off the stack, adding greedily to the matching.



S:




Select **any** edge e with positive weight. Reduce its weight from itself and its neighboring edges. Push e onto a stack. Repeat this procedure until no positive weight edges remain. Pop edges off the stack, adding greedily to the matching.



S:

Select **any** edge e with positive weight. Reduce its weight from itself and its neighboring edges. Push e onto a stack. Repeat this procedure until no positive weight edges remain. Pop edges off the stack, adding greedily to the matching.



Select **any** edge e with positive weight. Reduce its weight from itself and its neighboring edges. Push e onto a stack. Repeat this procedure until no positive weight edges remain. Pop edges off the stack, adding greedily to the matching.

Theorem ([Paz and Schwartzman, 2017])

The sequential local ratio algorithm gives a 2-approximation for weighted matching.

- ▶ Can cover a large fraction of the remaining sets with just a random sample of elements
- ▶ Random sample will be comparable to optimal solution since the processing order of local ratio is arbitrary

Initialize an empty stack S and repeat while E is not empty:

- ▶ Sample $8n^{1+\mu}$ edges uniformly from E
- ▶ Run local ratio on sampled edges, processing edges in order of **decreasing** weight, adding onto S
- ▶ Remove non-positive edges from E

Pop edges off S , adding greedily to the matching.

Initialize an empty stack S and repeat while E is not empty:


- ▶ Sample $8n^{1+\mu}$ edges uniformly from E
- ▶ Run local ratio on sampled edges, processing edges in order of **decreasing** weight, adding onto S
- ▶ Remove non-positive edges from E

Pop edges off S , adding greedily to the matching.


Theorem (Our paper)

Each iteration, the number of edges decrease by a factor of n^μ .


With $|E| = n^{1+c}$ edges, the entire algorithm terminates in $O(c/\mu)$ iterations.




MapReduce - Weighted Matching




MapReduce - Weighted Matching







- If the k -th heaviest edge adjacent to a node v is added to S , all but the top k edges are removed.



- Each node v has $O(n^\mu)$ edges sampled, so we get one of the top $O(\deg(v)/n^\mu)$ heaviest edges



- Degree of each node decreases by a factor of n^{μ} each iteration




- Further analysis gives $O(c/\mu)$ rounds instead of $O(1/\mu)$


Hungry-greedy

- ▶ Unlike local ratio, sampling is non-uniform
- ▶ **Heavy** elements are sampled first, to quickly disqualify a large fraction of elements from the solution


Given a graph $G = (V, E)$, an independent set (IS) is a set $S \subseteq V$ s.t. no two $v \in S$ are adjacent. A **maximal** independent set (MIS) is not a subset of any other IS.



Given a graph $G = (V, E)$, an independent set (IS) is a set $S \subseteq V$ s.t. no two $v \in S$ are adjacent. A **maximal** independent set (MIS) is not a subset of any other IS.



Given a graph $G = (V, E)$, an independent set (IS) is a set $S \subseteq V$ s.t. no two $v \in S$ are adjacent. A **maximal** independent set (MIS) is not a subset of any other IS.



Hungry-greedy - Maximal Independent Set



Hungry-greedy - Maximal Independent Set



Hungry-greedy - Maximal Independent Set





For $i = 1, \dots, 2/\mu$

- ▶ $V_H \leftarrow \{v \in V \mid \deg_I(v) > n^{1-i\mu/2}\}$
- ▶ While $|V_H| > n^{i\mu/2}$:
 - ▶ Draw $n^{i\mu/2}$ groups of $n^{\mu/2}$ vertices
 - ▶ For each group, add any vertices v to I if $\deg_I(v) > n^{i\mu/2}$
 - ▶ Remove any vertices from V_H with small degree
- ▶ Find an MIS in V_H and add to I

Theorem (Our paper)

Each iteration, the maximum degree of the graph decreases by a factor of $n^{\mu/4}$.

Problem	Approximation	MapReduce Rounds	Space per machine
Weighted b -matching	$3 - 2/b$	$O(c/\mu)$	$O(n^{1+\mu})$
Weighted vertex cover	2	$O(c/\mu)$	$O(n^{1+\mu})$
Weighted set cover	f	$O((c/\mu)^2)$	$O(f \cdot n^{1+\mu})$
	$(1 + \varepsilon) \ln \Delta$	$O\left(\frac{\log\left(\frac{w_{\max}}{w_{\min}} \Delta\right)}{\mu^2 \varepsilon}\right)$ if $n = \text{poly}(m)$	$O(m^{1+\mu})$
Maximal Clique		$O(1/\mu)$	$O(n^{1+\mu})$
Vertex Colouring	$(1 + o(1))\Delta$ colours	$O(1)$	$O(n^{1+\mu})$
Edge Colouring	$(1 + o(1))\Delta$ colours	$O(1)$	$O(n^{1+\mu})$

See full paper for definition of various parameters.

Techniques:

- ▶ Randomized local ratio
- ▶ Hungry-greedy

Thanks for listening! Questions?

The $\mu = 0$ case



Problem	Weighted?	Approximation	MapReduce Rounds	Space per machine
Vertex Cover	Y	2	$O(\log n)$	$O(n)$
Matching	Y	2	$O(\log n)$	$O(n)$
b -matching	Y	$3 - 2/b$	$O(\log n)$	$O(n)$
Maximal Indep. Set			$O(\log n)$	$O(n)$
Maximal Clique			$O(\log n)$	$O(n)$
Set Cover	Y	f	$O(\log^2 n)$	$O(f \cdot n)$
Vertex Colouring		$(1 + o(1))\Delta$ colours	$O(1)$	$\tilde{O}(n)$
Edge Colouring		$(1 + o(1))\Delta$ colours	$O(1)$	$\tilde{O}(n)$

Techniques:


- ▶ Randomized local ratio
- ▶ Hungry-greedy

Suppose $G = (V, E)$ with weights w_v on the vertices $v \in V$.
A minimum vertex cover is a set $S \subseteq V$ s.t. any $e \in E$ has at least one endpoint in S and the sum of weights in S is minimized.

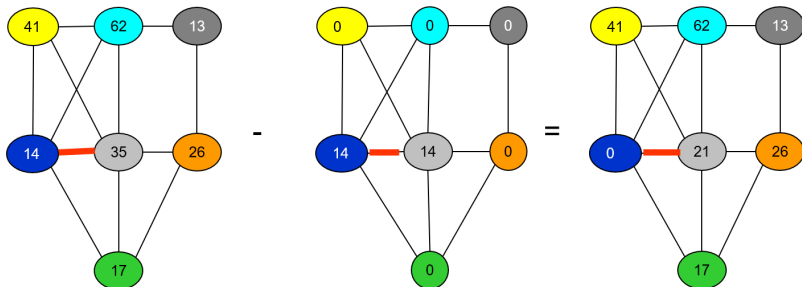
Sequential Local Ratio - Weighted Vertex Cover



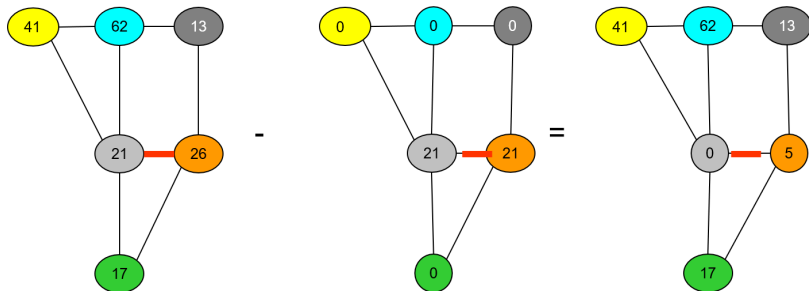
Select **any** uncovered edge e whose vertices have positive weight. Reduce the minimum of the two vertex weights from the two vertices. Remove all vertices with weight 0 and add them to the cover. Repeat while edges remain.



Select **any** uncovered edge e whose vertices have positive weight. Reduce the minimum of the two vertex weights from the two vertices. Remove all vertices with weight 0 and add them to the cover. Repeat while edges remain.



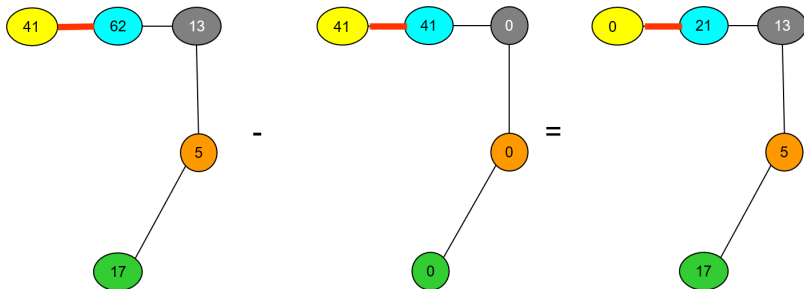
Select **any** uncovered edge e whose vertices have positive weight. Reduce the minimum of the two vertex weights from the two vertices. Remove all vertices with weight 0 and add them to the cover. Repeat while edges remain.



Sequential Local Ratio - Weighted Vertex Cover




Select **any** uncovered edge e whose vertices have positive weight. Reduce the minimum of the two vertex weights from the two vertices. Remove all vertices with weight 0 and add them to the cover. Repeat while edges remain.



Sequential Local Ratio - Weighted Vertex Cover




Select **any** uncovered edge e whose vertices have positive weight. Reduce the minimum of the two vertex weights from the two vertices. Remove all vertices with weight 0 and add them to the cover. Repeat while edges remain.



Sequential Local Ratio - Weighted Vertex Cover




Select **any** uncovered edge e whose vertices have positive weight. Reduce the minimum of the two vertex weights from the two vertices. Remove all vertices with weight 0 and add them to the cover. Repeat while edges remain.



Sequential Local Ratio - Weighted Vertex Cover



Select **any** uncovered edge e whose vertices have positive weight. Reduce the minimum of the two vertex weights from the two vertices. Remove all vertices with weight 0 and add them to the cover. Repeat while edges remain.



Select **any** uncovered edge e whose vertices have positive weight. Reduce the minimum of the two vertex weights from the two vertices. Remove all vertices with weight 0 and add them to the cover. Repeat while edges remain.

Theorem ([Bar-Yehuda and Even, 1985])

The sequential local ratio algorithm gives a 2-approximation for weighted vertex cover.

Repeat while E is not empty:

- ▶ Sample each edge in E with probability $2n^{1+\mu}/|E|$
- ▶ Run local ratio on sampled edges
- ▶ Add all vertices of zero weight to cover and update E

Theorem (Our paper)

Each iteration, the number of edges decrease by a factor of n^μ .

With $|E| = n^{1+c}$ edges, the entire algorithm terminates in $O(c/\mu)$ iterations.

Repeat while E is not empty:

- ▶ Sample $2n^{1+\mu}$ edges from E uniformly
- ▶ Run local ratio on sampled edges
- ▶ Add all vertices of zero weight to cover and update E

Intuition:

- ▶ Edges left for the next iteration must have both endpoints with positive weight.
- ▶ If there are more than $|E|/n^\mu$ such edges, then sampling $2n^{1+\mu}$ edges should have sampled all of them with high probability.
- ▶ If an edge is sampled, at least one of its endpoints must be zero weight, so it couldn't have been in the next iteration

Repeat while E is not empty:

- ▶ Sample $2n^{1+\mu}$ edges from E uniformly
- ▶ Run local ratio on sampled edges
- ▶ Add all vertices of zero weight to cover and update E

Can be further generalized to an f -approximation for set cover, where each element is in at most f sets.

- [Ahn and Guha, 2015] Ahn, K. J. and Guha, S. (2015).
Access to data and number of iterations: Dual primal algorithms for maximum matching under resource constraints.
In *Proceedings of the 27th ACM symposium on Parallelism in Algorithms and Architectures (SPAA)*, pages 202–211.
- [Bar-Yehuda and Even, 1985] Bar-Yehuda, R. and Even, S. (1985).
A local-ratio theorem for approximating the weighted vertex cover problem.
Annals of Discrete Mathematics, 25:27–45.
- [Crouch and Stubbs, 2014] Crouch, M. and Stubbs, D. S. (2014).
Improved streaming algorithms for weighted matching, via unweighted matching.
In *International Workshop on Approximation Algorithms for Combinatorial Optimization Problems (APPROX)*, pages 96–104.
- [Grigorescu et al., 2016] Grigorescu, E., Monemizadeh, M., and Zhou, S. (2016).
Streaming weighted matchings: Optimal meets greedy.
- [Karloff et al., 2010] Karloff, H. J., Suri, S., and Vassilvitskii, S. (2010).
A model of computation for MapReduce.
In *Proceedings of the Twenty-First Annual ACM-SIAM Symposium on Discrete Algorithms (SODA)*, pages 938–948.
- [Lattanzi et al., 2011] Lattanzi, S., Moseley, B., Suri, S., and Vassilvitskii, S. (2011).
Filtering: a method for solving graph problems in MapReduce.
In *Proceedings of the 23rd Annual ACM Symposium on Parallelism in Algorithms and Architectures (SPAA)*, pages 85–94.

- [[Leskovec et al., 2005](#)] Leskovec, J., Kleinberg, J. M., and Faloutsos, C. (2005).
Graphs over time: densification laws, shrinking diameters and possible explanations.
In *Proceedings of the Eleventh ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, Chicago, Illinois, USA, August 21-24, 2005*, pages 177–187.
- [[Paz and Schwartzman, 2017](#)] Paz, A. and Schwartzman, G. (2017).
A $(2 + \epsilon)$ -approximation for maximum weight matching in the semi-streaming model.
In *Proceedings of the Twenty-Eighth Annual ACM-SIAM Symposium on Discrete Algorithms (SODA)*, pages 2153–2161.
- [[Vassilvitskii, 2012](#)] Vassilvitskii, S. (2012).
CSCI 8980: Algorithmic techniques for big data analysis.